

```
In [3]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np

%matplotlib inline

In [5]: data = pd.read_csv("desktop/data/diabetes.csv")

In [6]: data.shape

Out[6]: (768, 9)

In [7]: data.head(5)

Out[7]:
   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  DiabetesPedigreeFunction  Age  Outcome
0           6     148             72           35         0  33.6                0.627    50         1
1           1      85             66           29         0  26.6                0.351    31         0
2           8     183             64           0          0  23.3                0.672    32         1
3           1      89             66           23        94  28.1                0.167    21         0
4           0     137             40           35        168  43.1                2.288    33         1

In [10]: # checking if any null value is present
data.isnull().values.any()

Out[10]: False

In [15]: import seaborn as sns
import matplotlib.pyplot as plt
corrmat = data.corr()
top_corr_features = corrmat.index
plt.figure(figsize=(20,20))
# plotting heatmap
g=sns.heatmap(data[top_corr_features].corr(),annot=True,cmap="RdYlGn")

In [16]: data.corr()

Out[16]:
   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  DiabetesPedigreeFunction  Age  Outcome
Pregnancies    1.000000    0.129459    0.141282   -0.081672  -0.073535    0.017683   -0.033523    0.544341    0.221898
Glucose         0.129459    1.000000    0.152590    0.057328    0.331357    0.221071    0.137337    0.263514    0.466581
BloodPressure   0.141282    0.152590    1.000000    0.207371    0.088933    0.281805    0.041265    0.239528    0.065068
SkinThickness  -0.081672    0.057328    0.207371    1.000000    0.436783    0.392573    0.183928   -0.113970    0.074752
Insulin        -0.073535    0.331357    0.088933    0.436783    1.000000    0.197859    0.185071   -0.042163    0.130548
BMI             0.017683    0.221071    0.281805    0.392573    0.197859    1.000000    0.140647    0.036242    0.292695
DiabetesPedigreeFunction -0.033523    0.137337    0.041265    0.183928    0.185071    0.140647    1.000000    0.033561    0.173844
Age             0.544341    0.263514    0.239528   -0.113970   -0.042163    0.036242    0.033561    1.000000    0.238356
Outcome         0.221898    0.466581    0.065068    0.074752    0.130548    0.292695    0.173844    0.238356    1.000000

In [25]: diabetes_true_count = len(data.loc[data['Outcome'] == True])
diabetes_false_count = len(data.loc[data['Outcome'] == False])

In [26]: (diabetes_true_count,diabetes_false_count)

Out[26]: (268, 500)

In [27]: #training test split
from sklearn.model_selection import train_test_split
features_columns = ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age']
predicted_class = ['Outcome']

In [47]: X = data[features_columns].values
y = data[predicted_class].values

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.30, random_state=10)

In [48]: print("total number of rows : {}".format(len(data)))
print("number of rows missing Glucose: {}".format(len(data.loc[data['Glucose'] == 0])))
print("number of rows missing BloodPressure: {}".format(len(data.loc[data['BloodPressure'] == 0])))
print("number of rows missing SkinThickness: {}".format(len(data.loc[data['SkinThickness'] == 0])))
print("number of rows missing insulin: {}".format(len(data.loc[data['Insulin'] == 0])))
print("number of rows missing BMI: {}".format(len(data.loc[data['BMI'] == 0])))
print("number of rows missing DiabetesPedigreeFunction: {}".format(len(data.loc[data['DiabetesPedigreeFunction'] == 0])))
print("number of rows missing Age: {}".format(len(data.loc[data['Age'] == 0])))

total number of rows : 768
number of rows missing Glucose: 5
number of rows missing BloodPressure: 35
number of rows missing SkinThickness: 227
number of rows missing insulin: 374
number of rows missing BMI: 11
number of rows missing DiabetesPedigreeFunction: 0
number of rows missing Age: 0

In [49]: from sklearn.impute import SimpleImputer
fill_values = SimpleImputer(missing_values=0, strategy="mean")
X_train = fill_values.fit_transform(X_train)
X_test = fill_values.fit_transform(X_test)

In [53]: ## Apply Algorithm
from sklearn.ensemble import RandomForestClassifier
random_forest_model = RandomForestClassifier(random_state=10)

random_forest_model.fit(X_train, y_train.ravel())

Out[53]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                               criterion='gini', max_depth=None, max_features='auto',
                               max_leaf_nodes=None, max_samples=None,
                               min_impurity_decrease=0.0, min_impurity_split=None,
                               min_samples_leaf=1, min_samples_split=2,
                               min_weight_fraction_leaf=0.0, n_estimators=100,
                               n_jobs=None, oob_score=False, random_state=10, verbose=0,
                               warm_start=False)

In [58]: predict_train_data = random_forest_model.predict(X_test)

from sklearn import metrics

print("Accuracy = {:.3f}").format(metrics.accuracy_score(y_test,predict_train_data)))

Accuracy = 0.766)
```